Shriram M

Namakkal, India | +91 96008 35006 | shriram.coder@gmail.com
linkedin.com/in/shriram-m-sde | github.com/shriram1206 | Personal Portfolio

Education

Selvam College of Technology — Namakkal, TN
Bachelor of Engineering in Computer Science and Engineering | ***CGPA: 8.0/10.0*** — *2022 – May 2026*

Experience

Software Development Engineer Intern — Feb 2026 – Mar 2026
Servimos Technologies Pvt Ltd (WeServeTech LLC) — *Chennai, TN*

- **Automated weekly compliance reporting** by building an AI report generator in Python and FastAPI, cutting manual preparation time from approximately 4 hours to under 10 minutes per cycle.
- **Implemented RAG-based context injection** via pgvector + PostgreSQL for semantic document retrieval and applied MCP (SSE transport, JSON-RPC tool calling) to automate recurring reporting pipelines.
- **Delivered a production-ready AI backend** over 3 Agile sprints with end-to-end ownership from schema design through deployment, with code and architecture reviewed and approved by senior engineers.
- **Added per-endpoint latency and error-rate visibility** to the reporting service using structured logging and custom middleware metrics, enabling rapid incident debugging after go-live.

AI/ML Intern — Jan 2025 – Feb 2025
YBI Foundation — *Remote*

- **Built and evaluated three ML models** (Random Forest, SVM, and Neural Network), achieving up to 89% validation accuracy on 10,000+ data points.
- **Implemented automated data cleaning and feature engineering pipelines** using Pandas and Scikit-Learn, reducing manual preprocessing effort by 70%.
- **Documented comprehensive experiments, metrics, and reproducible pipelines** with an emphasis on validation methodology.

Projects

Prompt Forge – AI Meta-Framework Synthesis Engine | *Python, FastAPI, Supabase, Groq* — 2026 | Live

- **Designed a full-stack SaaS compiler** that transforms vague user inputs into structured, tiered XML prompts by dynamically evaluating input complexity and applying a multi-stage cognitive framework pipeline.
- **Integrated Groq's Llama 3 API** via a Python/FastAPI backend, using strict metaprompting and Best-of-N gating to reduce hallucinations and improve output consistency.
- **Implemented API rate limiting and tier-based usage tracking** on Supabase using RLS, JWT auth, and a custom Debounce API to enforce cost controls without cross-tenant data leakage.

Axiom – Multi-Agent RAG Boardroom | *Python, FastAPI, Llama 3.1, ChromaDB, React* — 2026 | Source

- **Built a multi-agent RAG system** that simulates a corporate C-suite, utilizing asyncio.gather for concurrent, sub-second orchestration of 3 distinct LLM personas via the Groq API.
- **Implemented tenant isolation** using ChromaDB metadata filtering, ensuring department-scoped knowledge retrieval without data leakage.
- **Designed a shared-memory boardroom architecture** where agents dynamically evaluate, negotiate, and push back on cross-department proposals based on conflicting system prompts.

PharmaGraph Knowledge Engine – NL2Cypher Graph DB Tool | *Python, FastAPI, Neo4j, Groq (Llama-3)* 2026 | Source

- **Integrated a Neo4j graph database** coupled with Llama-3 (Groq API) to execute zero-shot NL2Cypher queries, dynamically resolving complex pharmacological entities with high precision.
- **Developed a Graph RAG pipeline** to map medical ontology queries into executable Cypher graph traversals, avoiding the multi-join overhead of equivalent relational queries.
- **Built a full-stack React/Vite dashboard** to natively visualize the executed analytical nodes in real-time, bridging plain language medical inquiries with robust data representations.

Technical Skills

AI & Agentic Systems: Python, LangChain, AI-Assisted Development, RAG, Prompt Engineering, Llama 3
Backend & APIs: FastAPI, Node.js, REST/JSON-RPC, SSE Transport
Databases (Vector & SQL): ChromaDB, pgvector, PostgreSQL, Supabase, MySQL
Frontend & Cloud: React.js, Tailwind CSS, Docker, Vercel, Git